

Software Design and Architecture

Building Scalable and Maintainable Systems

Understanding the foundations of great software systems

Instructor: Mehdi Lotfinejad

Learning Objectives

By the end of this lesson, you will be able to:

1. **Distinguish** between software design and architecture
2. **Explain** key principles: separation of concerns, modularity, abstraction
3. **Identify** major architectural patterns (Layered, MVC, Microservices, Event-Driven)
4. **Apply** SOLID principles at scale
5. **Evaluate** architectural trade-offs
6. **Design** system architectures for specific requirements

The House vs. City Analogy

Building a House (Design)

- Room layouts
- Plumbing connections
- Electrical wiring
- Interior details
- Furniture placement

Detailed implementation

Both essential, different scales, different concerns

Planning a City (Architecture)

- Neighborhoods
- Transportation systems
- Utility distribution
- District interactions
- Zone planning

High-level structure

Why Architecture Matters

The Cost of Bad Architecture:

- System unable to scale from 1K to 1M users
- Years paying off technical debt
- Expensive to change fundamental decisions
- Team productivity grinds to halt

The Benefit of Good Architecture:

- System adapts to changing needs
- Easy to maintain and enhance
- Scales with business growth
- Team moves fast with confidence

Architecture vs. Design: The Distinction

Architecture (System Level)

Big decisions:

- Monolith vs. microservices
- Communication protocols
- Data storage strategy
- Security boundaries
- Deployment strategy

Expensive to change

Design (Component Level)

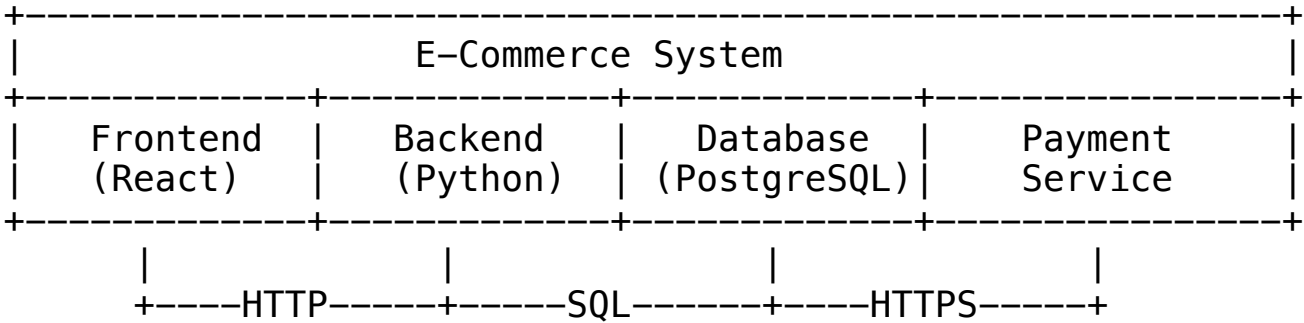
Implementation details:

- Class structures
- Algorithm choices
- Data structures
- Implementation patterns
- Code organization

Easier to change

Example: E-Commerce System

SYSTEM LEVEL (Architecture):



COMPONENT LEVEL (Design):

Backend Component Structure:

```
+-- controllers/
|   +-- ProductController (HTTP requests)
|   +-- OrderController (order processing)
|   +-- UserController (user management)
+-- services/
|   +-- ProductService (business logic)
|   +-- OrderService (order processing)
+-- repositories/
|   +-- ProductRepository (data access)
|   +-- OrderRepository (persistence)
+-- models/
|   +-- Product (data structures)
|   +-- Order (domain entities)
```

Architecture Determines Communication

Architecture Level:

- How components communicate: HTTP, gRPC, message queues
- Where data is stored: SQL, NoSQL, cache
- Security boundaries: authentication, authorization
- Deployment strategy: containers, VMs, serverless

Design Level:

- How classes interact within backend
- Which algorithms to use
- How to structure data
- Implementation patterns

Key Point: Architecture = System blueprint, Design = Component details

Fundamental Architectural Principles

Timeless principles that guide great architectures:

1. Separation of Concerns
2. Modularity and Encapsulation
3. Abstraction and Information Hiding
4. Loose Coupling and High Cohesion

These apply regardless of technology choices or programming languages

Principle 1: Separation of Concerns (SoC)

Core Idea:

Different aspects of the system should be isolated into distinct sections

A "concern" is a specific aspect:

- User interface
- Business logic
- Data access
- Logging
- Security
- Validation

Benefit: Change one aspect without touching others

SoC: Bad vs. Good Example

✗ Everything Mixed

```
def process_order(request):
    # Validation
    if not request.POST.get('id'):
        return "<html>Error</html>"

    # Database
    conn = psycopg2.connect(...)
    price = conn.execute(...)

    # Business logic
    total = price * 1.1

    # Persistence
    conn.execute("INSERT...")

    # Presentation
    return f"<html>{total}</html>"
```

✓ Clear Separation

```
# Presentation
class OrderController:
    def create(self, request):
        data = self._validate(request)
        order = self.service.create(data)
        return self._respond(order)

# Business Logic
class OrderService:
    def create(self, data):
        product = self.repo.find(data['id'])
        order = Order(product, data['qty'])
        return self.repo.save(order)

# Data Access
class OrderRepository:
    def save(self, order):
        # SQL operations
```

Benefits of Separation of Concerns

Change Database:

- Modify only Repository layer
- Business logic untouched
- UI remains unchanged

Change UI:

- Switch HTML to JSON
- Business logic unaffected
- Database queries unchanged

Change Business Rules:

- Update Service layer only
- UI and database unaffected

Principle 2: Modularity and Encapsulation

Modularity: Divide system into discrete, self-contained modules

Encapsulation: Hide internal details, expose only necessary interfaces

Think of kitchen appliances:

- Refrigerator has simple interface (door, controls)
- Internal complexity hidden (compressor, defrost)
- Can replace without rewiring kitchen

Modularity Example

```
class PaymentProcessor:
    """Public interface – stable contract"""

    def process_payment(self, amount, method, customer_id):
        """Client code depends on this stable interface"""
        self._validate_payment_request(amount, method)
        handler = self._get_payment_handler(method)
        result = self._process_with_retry(handler, amount, customer_id)
        self._log_transaction(result)
        return result

    # Private methods – hidden implementation
    def _validate_payment_request(self, amount, method): ...
    def _get_payment_handler(self, method): ...
    def _process_with_retry(self, handler, amount, customer_id): ...
    def _log_transaction(self, result): ...

# Client only sees public interface
processor = PaymentProcessor()
result = processor.process_payment(99.99, 'credit_card', 'CUST123')
```

Benefits of Modularity

Change Implementation:

- Add new payment methods
- Modify retry strategies
- Switch logging systems
- Update validation rules

Without breaking clients:

- Interface remains stable
- Client code unchanged
- Implementation evolves independently

Result: Flexibility and maintainability

Principle 3: Abstraction and Information Hiding

Abstraction: Focus on essential characteristics, hide unnecessary details

Information Hiding: Conceal implementation behind interfaces

Like driving a car:

- Interface: steering wheel, pedals, gear shift
- Hidden: engine timing, fuel injection, transmission ratios
- You don't need to understand internals to drive

Abstraction Example

```
from abc import ABC, abstractmethod

# Abstract interface – defines WHAT, not HOW
class DataStorage(ABC):
    @abstractmethod
    def save(self, key: str, data: dict) -> bool: pass

    @abstractmethod
    def retrieve(self, key: str) -> dict: pass

# Concrete implementations – HOW is hidden
class SQLStorage(DataStorage):
    def save(self, key, data):
        # SQL-specific implementation hidden
        self.conn.execute("INSERT INTO...", (key, json.dumps(data)))
        return True

class MongoStorage(DataStorage):
    def save(self, key, data):
        # MongoDB-specific implementation hidden
        self.collection.insert_one({'_id': key, 'data': data})
        return True
```


Using Abstraction

```
# Client depends on abstraction, not concrete class
class UserService:
    def __init__(self, storage: DataStorage): # Interface
        self.storage = storage

    def create_user(self, user_id, user_data):
        return self.storage.save(user_id, user_data)

# Can switch implementations without changing UserService
sql_storage = SQLStorage("postgresql://prod-db")
user_service = UserService(sql_storage)

# Later, switch to MongoDB – UserService unchanged
mongo_storage = MongoStorage("mongodb://new-db", "users")
user_service = UserService(mongo_storage)
```

Benefit: Change storage technology without touching business logic

Principle 4: Loose Coupling & High Cohesion

Loose Coupling

Minimal interdependence

Like separate appliances:

- Each plugged into outlet
- Independent operation
- Break one, others work
- Share only interface

High Cohesion

Elements strongly related

Like organized toolbox:

- Screwdriver drawer
- Hammer drawer
- Each clear purpose
- Related items together

Coupling: Bad vs. Good

✗ Tight Coupling

```
class OrderManager:
    def create_order(self, data):
        # Direct dependencies
        db = PostgreSQLDatabase()
        email = SMTPMailer()
        stripe = stripe.Client('key')

        # Locked to concrete classes
        # Hard to test
        # Hard to change
```

✓ Loose Coupling

```
class OrderService:
    def __init__(self,
                  repo,      # Interface
                  notifier,  # Interface
                  payment):  # Interface
        self.repo = repo
        self.notifier = notifier
        self.payment = payment

    def create_order(self, data):
        # Depends on abstractions
        # Easy to test (mock deps)
        # Easy to change (swap impl)
```

Benefits of Loose Coupling & High Cohesion

Loose Coupling:

- Change one part without breaking others
- Easy to test (mock dependencies)
- Reusable components
- Independent deployment

High Cohesion:

- Clear module purpose
- Easy to understand
- Changes localized
- Natural organization

Together: Maintainable, flexible systems

Major Architectural Patterns


Proven solutions to common structural problems

1. Layered (N-Tier) Architecture
2. Model-View-Controller (MVC)
3. Microservices Architecture
4. Event-Driven Architecture

Each pattern solves different problems at different scales

Pattern 1: Layered (N-Tier) Architecture

Organize code into horizontal layers



Layered Architecture: Code Example

```
# PRESENTATION LAYER
class ProductController:
    def __init__(self, product_service):
        self.product_service = product_service

    def get_products(self):
        category = request.args.get('category')
        products = self.product_service.find_products(category)
        return jsonify({'products': [self._format(p) for p in products]})

# BUSINESS LOGIC LAYER
class ProductService:
    def __init__(self, product_repo, inventory_repo):
        self.product_repo = product_repo
        self.inventory_repo = inventory_repo

    def find_products(self, category=None):
        products = self.product_repo.find_all()
        if category:
            products = [p for p in products if p.category == category]
        return [p for p in products if self.inventory_repo.is_in_stock(p.id)]
```

Layered Architecture: Data Layer

```
# DATA ACCESS LAYER
class ProductRepository:
    def __init__(self, db_connection):
        self.conn = db_connection

    def find_all(self):
        query = "SELECT id, name, price, category FROM products"
        cursor = self.conn.execute(query)
        return [self._row_to_product(row) for row in cursor.fetchall()]

    def _row_to_product(self, row):
        return Product(
            id=row[0],
            name=row[1],
            price=row[2],
            category=row[3]
        )
```


Layered Architecture: When to Use

✓ Benefits

- Clear separation
- Easy to understand
- Testable layers
- Team organization
- Structured approach

✓ Best For

- Enterprise apps
- CRUD operations
- Traditional monoliths
- Small to medium apps

✗ Drawbacks


- Performance overhead
- Can become rigid
- Database-centric
- Cross-layer changes

✗ Not Ideal For

- High-performance needs
- Rapidly changing requirements
- Complex business logic
- Microservices

Pattern 2: Model-View-Controller (MVC)

Separate concerns: data, presentation, and control



Controller handles input → updates **Model** → **View** displays

MVC: Model (Business Logic)

```
# MODEL: Business logic and data
class BlogPost:
    def __init__(self, title, content, author):
        self.title = title
        self.content = content
        self.author = author
        self.created_at = datetime.now()
        self.status = 'draft'

    def publish(self):
        if not self.title or not self.content:
            raise ValueError("Title and content required")
        self.status = 'published'
        self.published_at = datetime.now()

    def can_be_edited_by(self, user):
        return user.is_admin or user.username == self.author
```

MVC: View (Presentation)

```
# VIEW: Presentation logic
class BlogView:
    def render_post_list(self, posts):
        html = "<html><body><h1>Blog Posts</h1><ul>"
        for post in posts:
            html += f'<li><a href="/posts/{post.id}>">{post.title}</a></li>'
        html += "</ul></body></html>"
        return html

    def render_post_detail(self, post):
        html = f"""
        <html><body>
            <h1>{post.title}</h1>
            <p>By {post.author} on {post.created_at}</p>
            <div>{post.content}</div>
        </body></html>
        """
        return html

    def render_error(self, message):
        return f'<html><body><h1>Error</h1><p>{message}</p></body></html>'
```

MVC: Controller (Orchestration)

```
# CONTROLLER: Handles user input
class BlogController:
    def __init__(self, repository, view):
        self.repository = repository # Model
        self.view = view             # View

    def list_posts(self, request):
        posts = self.repository.find_all_published()
        html = self.view.render_post_list(posts)
        return Response(html, status=200)

    def update_post(self, request, post_id):
        post = self.repository.find_by_id(post_id)
        post.title = request.form['title']
        post.content = request.form['content']
        post.publish()
        self.repository.save(post)
        return Response(status=302, headers={'Location': f'/posts/{post.id}'})
```

MVC: When to Use

✓ Benefits

- Clear UI separation
- Parallel development
- Multiple views possible
- Testable business logic
- Framework support

✓ Best For

- Web applications
- Significant UI
- Multiple interfaces
- Django, Rails, ASP.NET

✗ Drawbacks

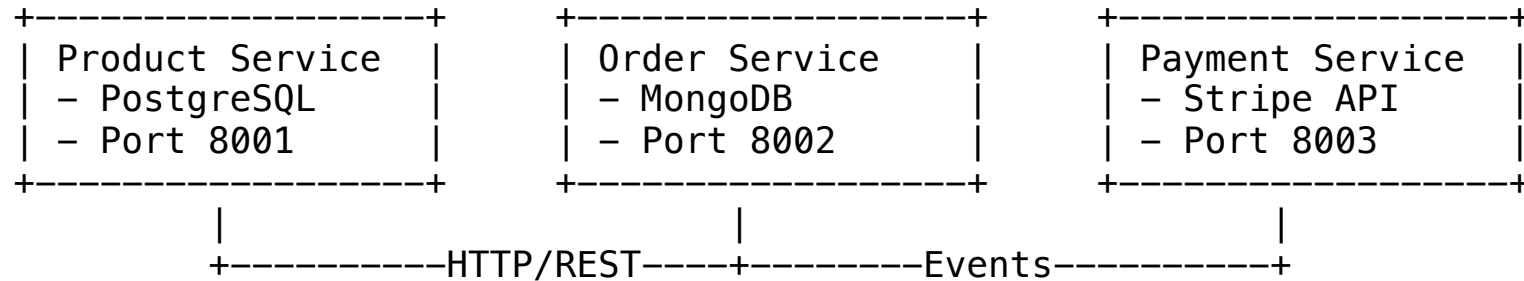
- Can be complex
- Responsibility confusion
- View-Controller coupling
- Overkill for simple apps

✗ Avoid When

- Simple CRUD
- No UI complexity
- API-only services
- Very small projects

Pattern 3: Microservices Architecture

Application as collection of small, independent services



Each service:

- Runs independently
- Own database
- Separate deployment
- Lightweight communication

Microservices: Product Service

```
# PRODUCT SERVICE (independent)
from flask import Flask, jsonify

app_products = Flask(__name__)

@app_products.route('/api/products/<product_id>', methods=['GET'])
def get_product(product_id):
    product = db.query_one(
        "SELECT * FROM products WHERE id = %s",
        (product_id,)
    )
    return jsonify({
        'product': {
            'id': product.id,
            'name': product.name,
            'price': product.price
        }
    })

if __name__ == '__main__':
    app_products.run(port=8001)
```


Microservices: Order Service

```
# ORDER SERVICE (calls Product Service)
import requests

@app_orders.route('/api/orders', methods=['POST'])
def create_order():
    data = request.get_json()

    # Call Product Service via HTTP
    products = []
    for item in data['items']:
        response = requests.get(
            f"http://product-service:8001/api/products/{item['id']}"
        )
        products.append(response.json()['product'])

    # Create order in this service's database
    order = {
        'items': data['items'],
        'total': sum(p['price'] * item['qty'] for p, item in zip(products, data['items']))
    }
    order_id = db.orders.insert_one(order).inserted_id

    # Publish event for other services
    event_bus.publish('order.created', {'order_id': str(order_id)})
    return jsonify({'order_id': str(order_id)}), 201
```

Microservices: Circuit Breaker Pattern

```
# Prevent cascading failures
class CircuitBreaker:
    def __init__(self, failure_threshold=5, timeout=60):
        self.failure_count = 0
        self.failure_threshold = failure_threshold
        self.timeout = timeout
        self.state = 'CLOSED' # CLOSED, OPEN, HALF_OPEN
        self.last_failure_time = None

    def call(self, func, *args, **kwargs):
        if self.state == 'OPEN':
            if time.time() - self.last_failure_time > self.timeout:
                self.state = 'HALF_OPEN'
            else:
                raise CircuitBreakerOpenError("Circuit breaker is open")

        try:
            result = func(*args, **kwargs)
            if self.state == 'HALF_OPEN':
                self.state = 'CLOSED'
                self.failure_count = 0
            return result
        except Exception as e:
            self.failure_count += 1
            self.last_failure_time = time.time()
            if self.failure_count >= self.failure_threshold:
                self.state = 'OPEN'
            raise
```

Microservices: When to Use

✓ Benefits

- Independent deployment
- Technology diversity
- Scalability per service
- Team autonomy
- Fault isolation

✓ Best For

- Large complex apps
- Multiple teams
- Different scaling needs
- Frequent updates
- DevOps maturity

✗ Drawbacks

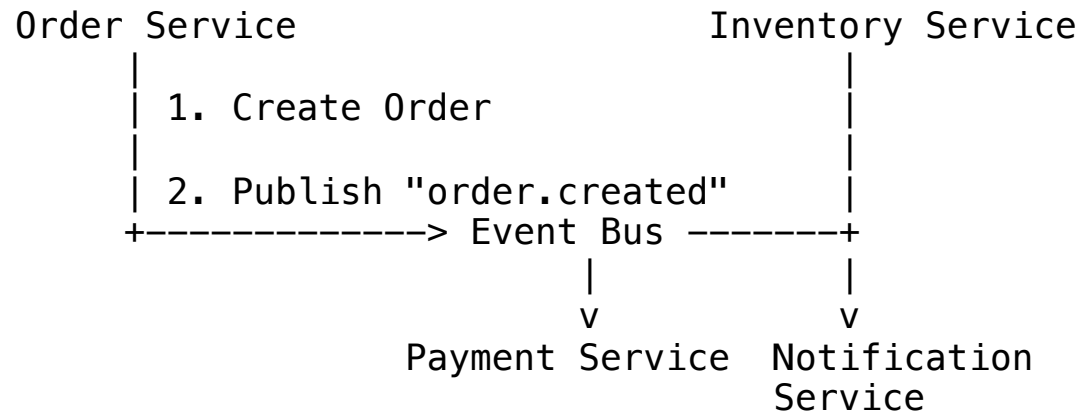
- Distributed complexity
- Network overhead
- Data consistency hard
- Testing difficulty
- Operational overhead

✗ Avoid When

- Small team (< 10 devs)
- Simple application
- No scaling needs
- Limited DevOps skills
- Tight coupling required

Pattern 4: Event-Driven Architecture

Components communicate via events



Asynchronous, decoupled, scalable

Event-Driven: Event Infrastructure

```
# Event Infrastructure
class Event:
    def __init__(self, event_type, data):
        self.event_type = event_type
        self.data = data
        self.event_id = str(uuid.uuid4())
        self.timestamp = datetime.now()

class EventBus:
    def __init__(self):
        self.subscribers = {}

    def publish(self, event):
        """Publish event - fire and forget"""
        if event.event_type in self.subscribers:
            for handler in self.subscribers[event.event_type]:
                # Asynchronous processing
                threading.Thread(target=handler, args=(event,)).start()

    def subscribe(self, event_type, handler):
        """Subscribe to specific event type"""
        if event_type not in self.subscribers:
            self.subscribers[event_type] = []
        self.subscribers[event_type].append(handler)
```

Event-Driven: Producer & Consumer

```
# EVENT PRODUCER
class OrderService:
    def __init__(self, event_bus):
        self.event_bus = event_bus

    def create_order(self, customer_id, items):
        order = {'order_id': str(uuid.uuid4()), 'customer_id': customer_id, 'items': items}
        self.repository.save(order)

        # Publish event - don't wait for consumers
        event = Event('order.created', order)
        self.event_bus.publish(event)
        return order['order_id']

# EVENT CONSUMER
class InventoryService:
    def __init__(self, event_bus):
        event_bus.subscribe('order.created', self.handle_order_created)

    def handle_order_created(self, event):
        order = event.data
        for item in order['items']:
            self.reserve_stock(item['product_id'], item['quantity'])
```

Event-Driven: When to Use

✓ Benefits

- Loose coupling
- Easy to scale
- Resilience
- Audit trail
- Add consumers easily

✓ Best For

- Complex workflows
- Multiple services
- Real-time updates
- Audit requirements
- Independent scaling

✗ Drawbacks

- Complexity
- Debugging difficulty
- Eventual consistency
- Schema management
- Learning curve

✗ Avoid When

- Simple workflows
- Immediate consistency needed
- Small team
- Synchronous requirements
- Tight deadlines

SOLID Principles at Scale

SOLID scales from classes to architectures

S - Single Responsibility Principle

O - Open/Closed Principle

L - Liskov Substitution Principle

I - Interface Segregation Principle

D - Dependency Inversion Principle

SOLID: Single Responsibility Principle

Each service/module should have one reason to change

✗ God Service

```
class UserService:
    def register_user(self): ...
    def authenticate(self): ...
    def send_email(self): ...
    def process_payment(self): ...
    def generate_reports(self): ...
```

Too many responsibilities!

✓ Focused Services

```
class UserRegistrationService:
    # Only registration

class AuthenticationService:
    # Only authentication

class NotificationService:
    # Only notifications

class BillingService:
    # Only billing
```

One responsibility each

SOLID: Open/Closed Principle

Open for extension, closed for modification

```
# Plugin architecture – extend without modifying core
class PaymentProcessor:
    def __init__(self):
        self.handlers = {}

    def register_handler(self, payment_type, handler):
        """Extend by adding new handlers"""
        self.handlers[payment_type] = handler

    def process(self, payment_type, amount):
        """Core logic unchanged"""
        if payment_type not in self.handlers:
            raise ValueError(f"Unknown payment type: {payment_type}")
        return self.handlers[payment_type].process(amount)

# Add new payment method WITHOUT modifying PaymentProcessor
processor.register_handler('crypto', CryptoPaymentHandler())
processor.register_handler('paypal', PayPalHandler())
```

SOLID: Liskov Substitution Principle

Services implementing interface should be substitutable

```
# Any implementation should work without breaking system
```

```
class MessageQueue(ABC):  
    @abstractmethod  
    def publish(self, message): pass  
  
    @abstractmethod  
    def consume(self): pass
```

```
# Can swap implementations
```

```
class RabbitMQQueue(MessageQueue): ...  
class KafkaQueue(MessageQueue): ...  
class SQSQueue(MessageQueue): ...
```

```
# System works with any implementation
```

```
def process_orders(queue: MessageQueue):  
    message = queue.consume()  
    # Works regardless of concrete queue type
```

SOLID: Interface Segregation Principle

Design focused, minimal APIs

✗ Fat Interface

```
class StorageService:
    def read_file(self): ...
    def write_file(self): ...
    def delete_file(self): ...
    def backup_database(self): ...
    def replicate_data(self): ...
    def compress_archive(self): ...
```

Too much in one interface

✓ Segregated

```
class FileStorage:
    def read(self): ...
    def write(self): ...

class DatabaseBackup:
    def backup(self): ...
    def restore(self): ...

class DataReplication:
    def replicate(self): ...
```

Focused interfaces

SOLID: Dependency Inversion Principle

Depend on abstractions, not concretions

```
# HIGH-LEVEL module depends on ABSTRACTION
class OrderService:
    def __init__(self, repository: OrderRepository): # Interface, not class
        self.repository = repository

    def create_order(self, data):
        order = Order(data)
        return self.repository.save(order)

# Can swap implementations without changing OrderService
sql_repo = SQLOrderRepository()
mongo_repo = MongoOrderRepository()
redis_repo = RedisOrderRepository()

# All work because they implement OrderRepository interface
order_service = OrderService(sql_repo) # Works
order_service = OrderService(mongo_repo) # Works
order_service = OrderService(redis_repo) # Works
```

Quality Attributes & Trade-offs

Architecture is fundamentally about trade-offs

Key Quality Attributes:

- **Scalability:** Handle increased load
- **Performance:** Response time, throughput
- **Availability:** Uptime percentage
- **Maintainability:** Ease of changes
- **Security:** Protection from threats
- **Testability:** Ease of testing
- **Deployability:** Ease of deployment

Every decision optimizes some while compromising others

Common Trade-offs: CAP Theorem

Consistency vs. Availability

You can't have both in distributed systems


Choose Consistency:

- Banking systems
- Financial transactions
- Inventory management

Choose Availability:

- Social media feeds
- Content delivery
- Shopping catalogs

CAP Triangle



Pick 2 of 3:

- CA: Not realistic (networks fail)
- CP: Consistent, might be unavailable
- AP: Available, eventually consistent

Common Trade-offs: Performance vs. Security

Performance vs. Security:

- Encryption adds latency
- Authentication adds overhead
- Authorization checks slow responses
- Security auditing impacts performance

Decision depends on data sensitivity:

- Medical records → Security wins
- Public blog → Performance wins
- Financial data → Security wins
- Marketing site → Performance wins

Common Trade-offs: Flexibility vs. Simplicity

More Flexible

- Plugin architecture
- Configuration-driven
- Multiple databases
- Extensible APIs

Cost:

- More complexity
- Harder to understand
- More code to maintain

Guideline: Start simple, add flexibility when needed (YAGNI)

More Simple

- Hardcoded choices
- Single database
- Fixed workflow
- Simple APIs

Cost:

- Less flexibility
- Changes require code
- Harder to extend

Common Trade-offs: Scalability vs. Consistency

```
# Trade-off example: Caching
class ProductService:
    def get_product(self, product_id):
        # Check cache first (fast but might be stale)
        cached = redis.get(f'product:{product_id}')
        if cached:
            return cached # FAST but potentially INCONSISTENT

        # Cache miss - hit database (slow but consistent)
        product = db.query("SELECT * FROM products WHERE id = %s", product_id)

        # Cache for 5 minutes
        redis.setex(f'product:{product_id}', 300, product)
        return product

# Trade-off: Performance + Scalability vs. Real-time Consistency
# Acceptable for: Product catalogs (prices don't change every second)
# Not acceptable for: Stock availability (need real-time inventory)
```

Documenting Decisions: ADRs

Architectural Decision Records capture the "why"

Title: Use Microservices for Order Processing

Status: Accepted

Date: 2024-01-15

Context:

- Monolith difficult to scale
- Order processing has 10x traffic of other components
- Different teams need independence

Decision:

Split order processing into separate microservice

Consequences:

- + Independent scaling of order component
- + Team can deploy without coordinating
- Added network latency for inter-service calls
- Need service discovery and monitoring
- Data consistency becomes harder

Alternatives Considered:

1. Keep monolith, optimize database
2. Use vertical scaling
3. Modular monolith approach



Practical Example: Social Media Platform

Let's design a scalable system from scratch

Business Requirements:

- 100M users
- 500M posts daily
- Real-time feed updates
- Follow/unfollow users
- Like, comment, share posts
- 99.9% availability required

How do we architect this?

Social Media: Service Decomposition

MICROSERVICES BY BUSINESS CAPABILITY:

- +-- User Service
 - | +-- Authentication
 - | +-- User profiles
 - | +-- Account management
- +-- Post Service
 - | +-- Create, edit, delete posts
 - | +-- Post metadata
- +-- Feed Service
 - | +-- Generate personalized feeds
 - | +-- Feed ranking algorithm
- +-- Social Graph Service
 - | +-- Follows, connections
 - | +-- Relationship queries
- +-- Engagement Service
 - | +-- Likes, comments, shares
 - | +-- Engagement counting
- +-- Media Service
 - | +-- Image/video upload
 - | +-- Storage management
- +-- Notification Service
 - | +-- Real-time alerts
 - | +-- Push notifications

Social Media: Communication Patterns

Synchronous (REST)

For immediate needs:

```
# Get user profile
user = user_service.get_user(
    user_id
)

# Upload image
url = media_service.upload(
    image_data
)
```

When: Need immediate response

Asynchronous (Events)

For scalability:

```
# Post created
event_bus.publish(
    "post.created",
    post_data
)

# Feed service consumes
# Notification service consumes
# Analytics service consumes
```

When: Don't need immediate response

Social Media: Data Storage Strategy

POLYGLOT PERSISTENCE – Right database for right job:

User Service	→ PostgreSQL (ACID for accounts, transactions)
Post Service	→ MongoDB (Flexible schema, document storage)
Feed Service	→ Redis (In-memory, ultra-fast reads)
Social Graph	→ Neo4j (Graph database for relationships)
Media Service	→ Amazon S3 + CloudFront CDN (Object storage for images/videos)
Analytics Service	→ Apache Cassandra (Time-series data, write-heavy)
Search Service	→ Elasticsearch (Full-text search)

Social Media: Feed Generation Strategy

Fan-out on Write

Push posts immediately

User posts ->
Push to all followers' feeds

Pros:

- + Fast reads
- + Precomputed feeds

Cons:

- High write cost
- Celebrities problem

For: Regular users

Solution: Hybrid approach based on user patterns

Fan-out on Read

Compute on demand

User requests feed ->
Query posts from followed users
Rank and return

Pros:

- + Low storage
- + Works for celebrities

Cons:

- Slow reads
- High compute

For: Inactive users, celebrities

Social Media: Scalability Strategy

```
# Hybrid Feed Generation
class FeedService:
    def generate_feed(self, user_id):
        user = self.user_service.get_user(user_id)

        if user.is_celebrity():
            # Fan-out on read (too many followers)
            return self._compute_feed_on_demand(user_id)

        if user.is_active():
            # Fan-out on write (precomputed)
            return self._get_precomputed_feed(user_id)

        # Inactive users: compute on demand
        return self._compute_feed_on_demand(user_id)

    def _get_precomputed_feed(self, user_id):
        # Fast: just read from Redis
        return redis.lrange(f'feed:{user_id}', 0, 50)

    def _compute_feed_on_demand(self, user_id):
        # Slower: query posts from followed users
        following = self.social_graph.get_following(user_id)
        posts = self.post_service.get_recent_posts(following)
        return self._rank_posts(posts)
```

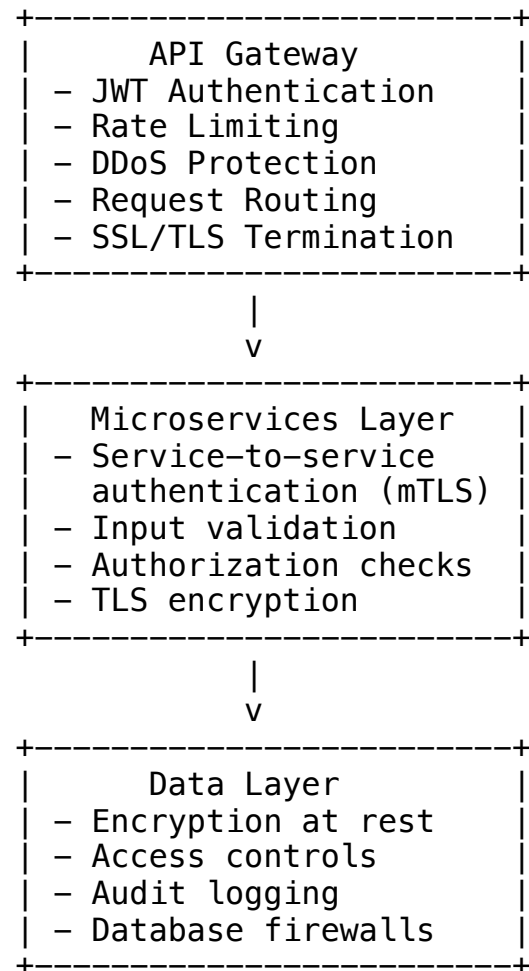
Social Media: Availability & Resilience

RESILIENCE PATTERNS:

- +-- Multi-Availability Zone Deployment
 - | +-- Services across 3+ zones
 - | +-- Load balancing
 - |
- +-- Circuit Breakers
 - | +-- Prevent cascading failures
 - | +-- Fail fast, don't wait
 - |
- +-- Graceful Degradation
 - | +-- Show cached feeds if real-time unavailable
 - | +-- Disable non-critical features
 - |
- +-- Retry with Backoff
 - | +-- Retry failed requests
 - | +-- Exponential backoff
 - |
- +-- Rate Limiting
 - | +-- Prevent overload
 - | +-- Per-user limits
 - |
- +-- Database Replication
 - +-- Master-slave replication
 - +-- Read replicas for scaling

Social Media: Security Architecture

SECURITY LAYERS:



Social Media: Evolution Path

PHASE 1: MVP (0–10K users)

- Monolithic application
- Single server
- PostgreSQL database
- Simple caching

PHASE 2: Growth (10K–1M users)

- Modular monolith
- Separate frontend/backend
- Read replicas
- Redis caching
- CDN for static assets

PHASE 3: Scale (1M–10M users)

- Microservices for high-traffic
- Message queue
- Multiple database types
- Container orchestration (Kubernetes)

PHASE 4: Hypergrowth (10M+ users)

- Full microservices
- Event-driven architecture
- Global CDN
- Multi-region deployment
- Advanced caching strategies

Common Pitfalls

Pitfall #1: Premature Optimization

Problem:

Starting with microservices for 100 users

Example:

- 3 developers
- 100 users
- Simple CRUD app
- Choose microservices "because Netflix uses them"

Result:

- Operational nightmare
- Slow development

Pitfall #2: Distributed Monolith

Problem:

Microservices that are tightly coupled

Warning Signs:

- Deploying one service requires deploying three others
- Shared database between services
- Services calling each other synchronously in chains
- No clear service boundaries

Result:

Worst of both worlds:

- Complexity of microservices
- Coupling of monolith

Solution:

loftinejad.com

Clear service boundaries, async communication, independent data stores

Pitfall #3: Ignoring Conway's Law

Conway's Law:

"Organizations design systems that mirror their communication structure"

Example:

- Frontend team (10 people)
- Backend team (10 people)
- Database team (5 people)

Result: Frontend monolith + Backend monolith + Shared database

Better:

Organize teams around business capabilities:

- User Management team (owns full stack)
- Content team (owns full stack)
- Payments team (owns full stack)

Pitfall #4: No Clear Boundaries

Bad Service Design:

UserService:

- User registration
- Order processing
- Email sending
- Report generation

Too many responsibilities!

Good Service Design:

UserService: User management only
OrderService: Order processing only
NotificationService: Email/SMS only
ReportService: Analytics only

Clear, focused responsibilities

Best Practices

Best Practice #1: Start Simple, Evolve

EVOLUTION PATH:

Monolith -> Modular Monolith -> Microservices

WHEN TO SPLIT:

- + Performance bottleneck in specific component
- + Team too large for single codebase
- + Different scaling needs for components
- + Clear boundaries established

DON'T SPLIT WHEN:

- Team is small (< 10 developers)
- Requirements are unclear
- User base is small
- No operational complexity needed

Remember: Premature distribution = premature optimization

Best Practice #2: Design for Failure

In distributed systems, failures are normal

```
# Timeouts
response = requests.get(url, timeout=5)

# Retries with exponential backoff
@retry(tries=3, delay=1, backoff=2)
def call_external_service():
    return requests.get(url)

# Circuit breakers
breaker = CircuitBreaker(failure_threshold=5, timeout=60)
result = breaker.call(risky_operation)

# Graceful degradation
try:
    recommendations = recommendation_service.get_recommendations()
except ServiceUnavailable:
    recommendations = get_cached_recommendations()
    # System still works, just with stale data
```

Best Practice #3: Observability First

You can't fix what you can't see

Three Pillars of Observability:

1. Logging:

- Structured logs (JSON)
- Correlation IDs for tracing
- Centralized aggregation (ELK, Splunk)

2. Metrics:

- Response times, error rates
- Resource usage (CPU, memory)
- Business metrics (orders/min)

3. Tracing:

lotfinejad.com

- Distributed tracing (Jaeger, Zipkin)

Best Practice #4: API Versioning

Services will evolve. Plan for it.

```
# URL versioning
@app.route('/api/v1/users')
def get_users_v1():
    return jsonify(users)


@app.route('/api/v2/users')
def get_users_v2():
    # Enhanced with pagination
    return jsonify({
        'users': users,
        'page': page,
        'total': total
    })

# Deprecation strategy:
# v1: Current (supported 12 months)
# v2: New (encouraged)
# v3: Beta (testing)
# After 6 months: Announce v1 deprecation
# After 12 months: Remove v1
```

Best Practice #5: Document Architecture

Use C4 Model for clarity:

LEVEL 1: System Context



LEVEL 2: Containers



LEVEL 3: Components

API Gateway:

- +-- Authentication
- +-- Rate Limiting
- +-- Routing
- +-- Monitoring

LEVEL 4: Code

(Class diagrams, detailed design)



Practice Quiz #1

Healthcare Scenario:

Medical device software with FDA regulations. Requirements defined by regulatory standards. Extensive documentation needed for approval. Safety-critical with thorough testing before release.

Question:

Which architectural approach is MOST appropriate?

- A) Microservices with continuous deployment
- B) Event-driven architecture with message queues
- C) Layered monolithic with comprehensive documentation
- D) Serverless architecture with AWS Lambda

Think before next slide...

Quiz #1: Answer

Answer: C) Layered monolithic with comprehensive documentation

Why Layered Monolith:

- ✓ **Stable requirements** (defined by FDA)
- ✓ **Extensive documentation** (regulatory requirement)
- ✓ **Fixed compliance standards**
- ✓ **Safety-critical** (thorough testing before release)
- ✓ **Full audit trail**

Why Not Others:

- ✗ **Microservices:** Insufficient documentation, incremental releases not allowed
- ✗ **Event-driven:** Eventual consistency not acceptable for medical devices
- ✗ **Serverless:** Execution limits, cold starts, lack of control



Practice Quiz #2

E-Commerce Scenario:

Product catalog shows prices from inventory database. Every page load queries database directly (2-3 seconds during peak). Considering Redis cache with 5-minute TTL.

Question:

What trade-off are you making?

- A) Trading security for performance
- B) Trading consistency for performance
- C) Trading availability for scalability
- D) Trading maintainability for flexibility

Think before next slide...

Quiz #2: Answer

Answer: B) Trading consistency for performance

Explanation:

Consistency Sacrifice:

- Prices might be up to 5 minutes stale
- Database shows \$89, cache shows old \$99
- Users see inconsistent data temporarily

Performance Gain:

- Page loads: 2-3 seconds → milliseconds
- Database load dramatically reduced
- Can handle much more traffic

When Acceptable:





Practice Quiz #3

Scaling Scenario:

Which provides STRONGEST justification for microservices over monolith?

- A) Team has 3 developers, needs to launch quickly
- B) Different parts have vastly different scaling needs (1000x traffic difference)
- C) Want to use multiple programming languages
- D) System needs to be highly available

Think before next slide...

Quiz #3: Answer

Answer: B) Different parts have vastly different scaling needs

Why This Justifies Microservices:

Example: E-commerce platform

- **Search Service:** 100,000 requests/min (users browsing)
- **Checkout Service:** 100 requests/min (only 0.1% purchase)

With Monolith:

- Must scale entire app for search traffic
- Waste resources on rarely-used checkout

With Microservices:

- Deploy 50 search instances
- Deploy 2 checkout instances

- Massive cost savings



Practice Quiz #4

Code Structure Scenario:

Which BEST demonstrates separation of concerns?

- A) Single function: validates input, queries database, performs calculations, renders HTML
- B) Controller handles HTTP, Service contains business logic, Repository handles data
- C) All database queries in controller, business logic spread across layers
- D) Business logic in view templates, controllers only route requests

Think before next slide...

Quiz #4: Answer

Answer: B) Controller handles HTTP, Service contains business logic, Repository handles data

Why This Is Correct:

Controller Layer:

- HTTP concerns only
- Parse requests, format responses
- No business logic or database queries

Service Layer:

- Business logic only
- Calculations, validations, workflows
- No HTTP or SQL knowledge

Repository Layer:

- Data access only



Practice Quiz #5

Event-Driven Scenario:

Order Service publishes "order.created" event. Three subscribers: Inventory (reserves stock), Payment (charges customer), Notification (sends email). Payment service fails due to outage.

Question:

What happens?

- A) Entire order creation fails and rolls back
- B) Inventory and Notification succeed; Payment processes when recovered
- C) All three services fail (tightly coupled)
- D) System enters unrecoverable inconsistent state

Think before next slide...

Quiz #5: Answer

Answer: B) Inventory and Notification succeed; Payment processes when recovered

How Event-Driven Works:

1. **Order Service publishes** event to message broker
2. **Message broker stores** event persistently
3. **Inventory subscribes**, receives immediately, reserves stock ✓
4. **Notification subscribes**, receives immediately, sends email ✓
5. **Payment is down**, broker keeps event in Payment's queue
6. **When Payment recovers**, processes all queued events ✓

Key Benefits:

- ✓ **Resilience:** One service failure doesn't cascade
- ✓ **Decoupling:** Services don't know about each other
- ✓ **Eventual Consistency:** System becomes consistent eventually

Key Takeaways

Remember Forever:

1. **Architecture ≠ Design**: Architecture = system structure, Design = component implementation
2. **Principles guide decisions**: SoC, modularity, abstraction, loose coupling
3. **Patterns solve specific problems**: Layered, MVC, Microservices, Event-Driven
4. **Trade-offs are inevitable**: Optimize some attributes, compromise others
5. **SOLID scales to systems**: Not just classes, but entire architectures
6. **Start simple, evolve**: Monolith → Modular Monolith → Microservices
7. **Document decisions**: Use ADRs to capture context and rationale

Universal Success Factors

Regardless of architecture chosen:

- ✓ **Clear communication** between team and stakeholders
- ✓ **Skilled team** competent in chosen approach
- ✓ **Committed stakeholders** engaged appropriately
- ✓ **Quality focus** on testing and code quality
- ✓ **Realistic planning** with honest estimates
- ✓ **Risk management** identifying and mitigating risks
- ✓ **Continuous learning** from retrospectives
- ✓ **Leadership support** with executive sponsorship

Bottom Line: Well-executed architecture > poorly-executed "best practice"

Action Items

This Week:

1. Assess your current project's architecture
2. Identify one architectural principle being violated
3. Propose one specific improvement
4. Share learning with your team

This Month:

- Study one architectural pattern in depth
- Read one book from references
- Review architectural decisions in your codebase
- Hold architecture retrospective with team

Resources for Continued Learning

Essential Reading:

- **Domain-Driven Design** by Eric Evans
- **Building Microservices** by Sam Newman
- **Clean Architecture** by Robert Martin
- **Software Architecture in Practice** by Bass, Clements, Kazman
- **Fundamentals of Software Architecture** by Richards and Ford

Online Resources:

- Martin Fowler's Blog (martinfowler.com)
- Microsoft Architecture Center
- AWS Well-Architected Framework
- C4 Model (c4model.com)

Architecture Decision Framework

WHEN CHOOSING ARCHITECTURE:

1. Assess Requirements
 - Stability (stable → Waterfall, evolving → flexible)
 - Scale (small → monolith, large → distributed)
 - Team size (small → simple, large → microservices)
2. Identify Constraints
 - Budget, timeline, team skills
 - Regulatory, compliance requirements
 - Technology limitations
3. Evaluate Trade-offs
 - What are you optimizing for?
 - What are you willing to sacrifice?
 - Document in ADRs
4. Start Simple
 - Don't prematurely distribute
 - Establish clear boundaries
 - Split when scaling demands
5. Measure and Adapt
 - Monitor quality attributes
 - Gather feedback
 - Evolve architecture based on data

Final Thoughts

Software Architecture Is:

- About trade-offs, not perfection
- Context-dependent, not one-size-fits-all
- Evolutionary, not set in stone
- Balancing multiple concerns

Success Comes From:

- Understanding principles deeply
- Matching architecture to context
- Documenting decisions and rationale
- Evolving based on feedback
- Prioritizing quality attributes

Thank You!

Questions?

Key Message:

Architecture is about making informed trade-offs that balance competing quality attributes while keeping your system maintainable, scalable, and aligned with business goals.

Choose wisely, document thoroughly, evolve thoughtfully.

Lesson Complete

You Are Now Equipped To:

- ✓ Distinguish architecture from design
- ✓ Apply fundamental principles (SoC, modularity, abstraction)
- ✓ Choose appropriate architectural patterns
- ✓ Evaluate trade-offs between quality attributes
- ✓ Apply SOLID principles at system level
- ✓ Design for failure and resilience
- ✓ Document architectural decisions
- ✓ Evolve architecture thoughtfully

Now go build amazing systems!

Keep learning, keep building, keep improving!

